

Network simulation and TCP congestion control analysis using ns3

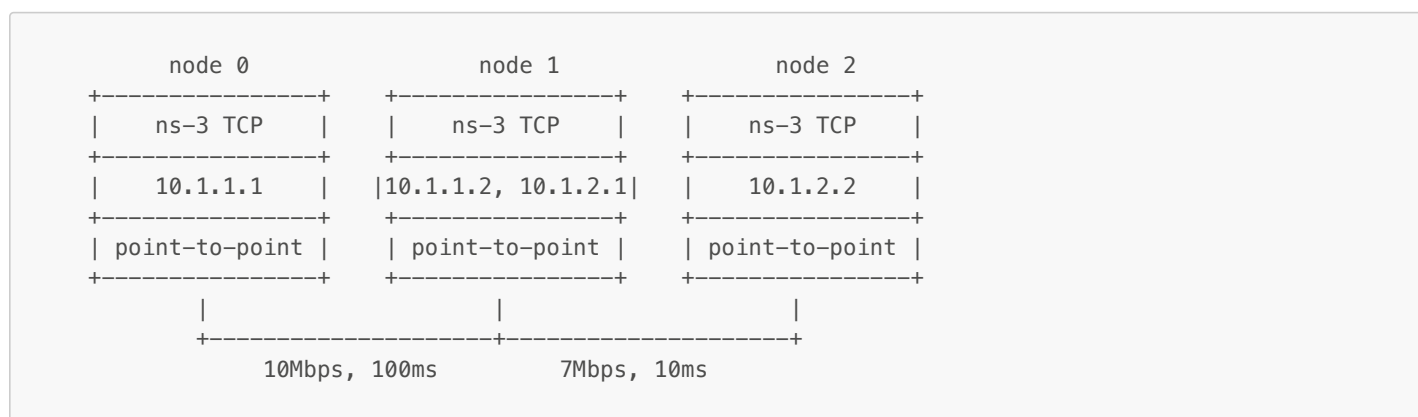
Saksham Singh

2022434

CSE 232 - Computer Networks - PA04

Network Topology

- Three nodes are connected in a linear topology. The first node **n0** is connected to the second node **n1** and the second node is connected to the third node **n3**.
- The link between the nodes is a point-to-point link, with node **n0** and **n1** connected with a link of **10Mbps** bandwidth and **100ms** delay. The link between node **n1** and **n2** is of **7Mbps** bandwidth and **10ms** delay.
- The queue size is set to **50** packets of size **1460** bytes, and error rate is set to **0.000001**. The TCP variant used is **TcpNewReno**, and the ns3 simulation time is set to **10s**.



Q1 - Run the simulation with the default parameters (provided in the table) and answer the following questions

a) What is the maximum expected value (theoretical) of throughput (in Mbps)? Why?

- The maximum theoretical throughput can be **7Mbps** because the bottleneck link is between node **n1** and **n2** which has a bandwidth of **7Mbps**. The throughput of the network is limited by the bandwidth of the bottleneck link.

b) How much is Bandwidth-Delay-Product (BDP)? Express your answer in terms of the number of packets.

- The Bandwidth-Delay-Product (BDP) is the maximum number of bits that can be in the network at any given time. It is calculated as the product of the bandwidth and the round-trip time (RTT) of the network. The RTT of the network is the sum of the delays of the links between the nodes. The BDP is calculated as:

$$\text{BDP} = \text{Bandwidth} * \text{RTT_Delay}$$

The RTT delay can be computed as the sum of propagation delays of the links between the nodes and the transmission delays of the packets. The propagation delay is the time taken for the packet to travel from the source to the destination, and the transmission delay is the time taken to transmit the packet over the link. The propagation delay for the two links is: **100ms + 10ms = 110ms**. The transmission delay can be calculated as the time taken to transmit a packet of size **1460** bytes over the link, defined as **Packet size / Maximum bandwidth = 1460 bytes / 7Mbps = 0.00166857 secs = 1.66857ms**.

Therefore, the transmission side of the RTT is **110ms + 1.66857ms = 111.66857ms**.

For the round trip, assuming the acknowledgement is sent immediately after the packet is received, and acknowledgement transmission delay is 0ms, the return trip time is just the propagation delay of the two links, which is **110ms**.

Therefore, the RTT is **111.66857ms + 110ms = 221.66857ms**.

- Calculating the Bandwidth-Delay-Product (BDP) as:

$$\text{BDP} = 7\text{Mbps} * 221.66857\text{ms} = 7 * 10^6 * 221.66857 * 10^{-3} = 1551680 \text{ bits} = 193960 \text{ bytes}$$

The packet size is 1460 bytes,

$$\text{BDP} = 193960 \text{ bytes} / 1460 \text{ bytes} = 132.8 \text{ packets}$$

Therefore, the Bandwidth-Delay-Product (BDP) is 132.8 packets.

- If we assume that acknowledgement is not sent, BDP will be $7\text{Mbps} * 111.66857\text{ms} = 776679 \text{ bits} = 97084 \text{ bytes} = 66.4 \text{ packets}$, roughly half of the previous value.

c) What is the average computed throughput of the TCP transfer?

- Upon running the simulation and analyzing the .pcap files for all the nodes, we can calculate the sum of the bytes received and divide by the simulation time to get the average throughput. It can be achieved by running the following command:

```
tshark -r {file} -q -z io,stat,10,"SUM(frame.len)frame.len"
```

where {file} is the .pcap file for the node. The output will be as follows:

```
=====
| I/O Statistics                                     |
| Duration: {duration} secs                         |
| Interval: {duration} secs                        |
| Col 1: SUM(frame.len)frame.len                   |
|-----|-----|
| Interval          | 1  SUM  |
|-----|-----|
| 0.000 <> {duration} | {bytes} |
|-----|-----|
=====
```

The average throughput can be calculated as $\text{bytes} / \text{duration}$.

- This can be automated using the following commands in python:

```
files = ["Q1_tcp-example-0-0.pcap", "Q1_tcp-example-1-0.pcap", "Q1_tcp-example-1-1.pcap",
"Q1_tcp-example-2-0.pcap"]
throughput = []

for file in files:
    # Parse the pcap file for throughput
    command = f'tshark -r {file} -q -z io,stat,10,\"SUM(frame.len)frame.len\"'
    output = os.popen(command).read()
    # print(output)
    output = output.split()
    # print(output)
    frame_len = float(output[-4])
    duration = float(output[-6])
    throughput.append(frame_len * 8 / duration / 1000000)
```

```
for i, file in enumerate(files):
    print(f"Throughput for {file} is {throughput[i]} Mbps")
```

- The following are the average throughputs for the nodes:

```
Throughput for Q1_tcp-example-0-0.pcap is 3.2745145403899723 Mbps
Throughput for Q1_tcp-example-1-0.pcap is 3.26832 Mbps
Throughput for Q1_tcp-example-1-1.pcap is 3.2656683146067413 Mbps
Throughput for Q1_tcp-example-2-0.pcap is 3.261553830577118 Mbps
```

d. Is the achieved throughput approximately equal to the maximum expected value? If it is not, explain the reason for the difference.

- As we can see, the theoretical maximum throughput is **7Mbps**, but the achieved throughput is around **3.26Mbps**. The achieved throughput is less than the theoretical maximum throughput because of the following reasons:
 - The TCP congestion control algorithm used is **TcpNewReno**, which is a conservative algorithm that reduces the congestion window size by half upon detecting congestion. This reduces the throughput of the network.
 - The network has a queue size of **50** packets, which can cause packet loss due to buffer overflow. Packet loss can cause the TCP sender to reduce the congestion window size, which reduces the throughput of the network.
 - The error rate of the network is set to **0.000001**, which can cause packet loss due to bit errors. Packet loss can cause the TCP sender to reduce the congestion window size, which reduces the throughput of the network.
- These factors combined can reduce the throughput of the network from the theoretical maximum value.

e. Plot Congestion Window (CWND) with time


- The congestion window (CNWD) can be plotted by parsing the **.cwnd** file generated using the ns3 simulation, and plotted using matplotlib. It can be done using the following python code:

```
import matplotlib.pyplot as plt

with open("Q1_tcp-example.cwnd", "r") as f:
    time = []
    cwnd_low = []
    cwnd_high = []
    for line in f:
        fields = line.strip().split()
        time.append(float(fields[0]))
        cwnd_low.append(float(fields[1]))
        cwnd_high.append(float(fields[2]))

# Plot the cwnd
plt.figure(figsize=(20, 6))
plt.plot(time, cwnd_low, label="cwnd_low")
plt.plot(time, cwnd_high, label="cwnd_high")
plt.xlabel("Time (s)")
plt.ylabel("cwnd")
plt.title("Q1 - cwnd vs Time")
plt.legend()
plt.grid()
plt.savefig("Q1_cwnd.png")
plt.show()
```

- The plot will be saved as **Q1_cwnd.png** and will look like the following:



f. Plot queueing delay with time

- The queueing delay can be plotted by parsing the trace file generated using the ns3 simulation, and plotted using matplotlib. It can be done using the following python code:

```
import matplotlib.pyplot as plt

time = []
queuing_delay = []
enqueued_times = {}

# Parse the trace file for queueing delay -- format as below
# + time
# - time
# r time

for line in open(trace_file):
    fields = line.strip().split()
    event = fields[0]
    timestamp = float(fields[1])
    node_id = fields[2]
    packet_info = " ".join(fields[3:])


    if event == "+":
        key = packet_info
        enqueued_times[key] = timestamp
    elif event == "-":
        key = packet_info
        if key in enqueued_times:
            delay = timestamp - enqueued_times[key]
            time.append(timestamp)
            queuing_delay.append(delay)
            del enqueued_times[key]

# save the data to a file
with open("Q1_QueueingDelay.txt", "w") as f:
    for i in range(len(time)):
        f.write(f"{time[i]} {queuing_delay[i]}\n")

# Plot the queueing delay
plt.figure(figsize=(20, 6))
plt.plot(time, queuing_delay)
plt.xlabel("Time (s)")
plt.ylabel("Queueing Delay (s)")
plt.title("Q1 - Queueing Delay vs Time")
plt.grid()
```

```
plt.savefig("Q1_QueueingDelay.png")
plt.show()
```

- The plot will be saved as `Q1_QueueingDelay.png` and will look like the following:



- Similarly, the queuing length can be plotted using the following python code:

```
import matplotlib.pyplot as plt

time = []
queuing_length = []
enqueued_packets = {}

# Parse the trace file for queuing length -- format as below
# + time
# - time
# r time

for line in open(trace_file):
    fields = line.strip().split()
    event = fields[0]
    timestamp = float(fields[1])
    node_id = fields[2]
    packet_info = " ".join(fields[3:])

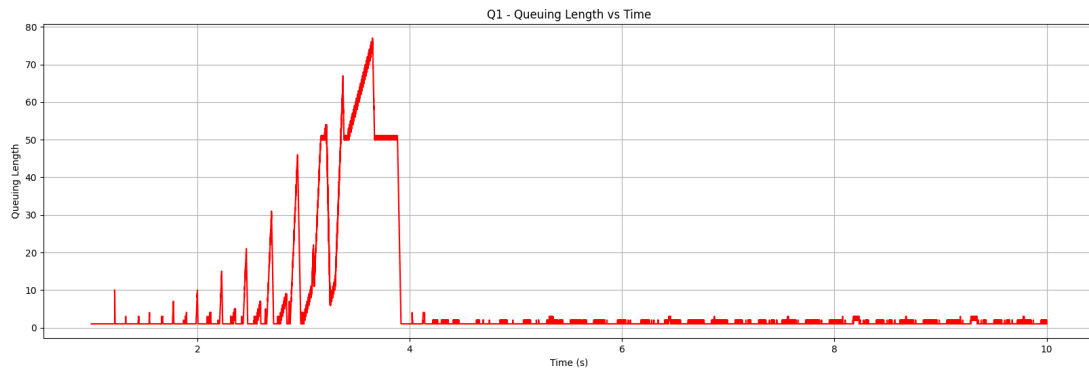
    if event == "+":
        key = packet_info
        enqueued_packets[key] = timestamp
    elif event == "-":
        key = packet_info
        if key in enqueued_packets:
            time.append(timestamp)
            queuing_length.append(len(enqueued_packets))
            del enqueued_packets[key]

# save the data to a file
with open("Q1_QueueingLength.txt", "w") as f:
    for i in range(len(time)):
        f.write(f"{time[i]} {queuing_length[i]}\n")

# Plot the queuing delay
plt.figure(figsize=(20, 6))
plt.plot(time, queuing_length)
plt.xlabel("Time (s)")
plt.ylabel("Queueing Length")
plt.title("Q1 - Queueing Length vs Time")
plt.grid()
```

```
plt.savefig("Q1_QueueingLength.png")
plt.show()
```

- The plot will be saved as **Q1_QueueingLength.png** and will look like the following:



g) Are the plots in 1(e) and 1(f) related? Explain the relationship between the two plots.

- The plots in 1(e) and 1(f) are related because the queueing delay is related to the congestion window size. When the congestion window size is small, the packets are sent slowly, which reduces the queueing delay. When the congestion window size is large, the packets are sent quickly, which increases the queueing delay. Therefore, the queueing delay is proportional to the congestion window size, and the two plots are related.

Q2 - Change queue size to 1000 (rest of the parameter values are same as default values)

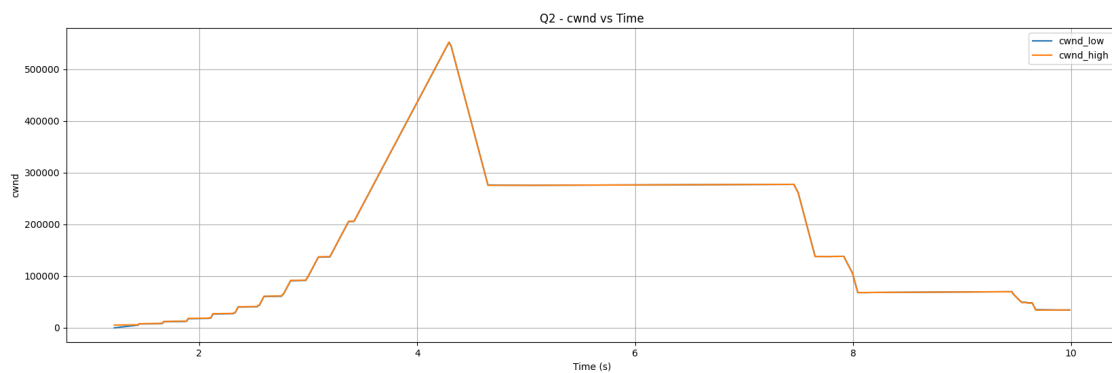
a) What is the average computed throughput of the TCP transfer?

- The average throughput values are:

```
Throughput for Q2_tcp-example-0-0.pcap is 4.808675642594859 Mbps
Throughput for Q2_tcp-example-1-0.pcap is 4.833663108214406 Mbps
Throughput for Q2_tcp-example-1-1.pcap is 4.833663108214406 Mbps
Throughput for Q2_tcp-example-2-0.pcap is 4.835047361907977 Mbps
```

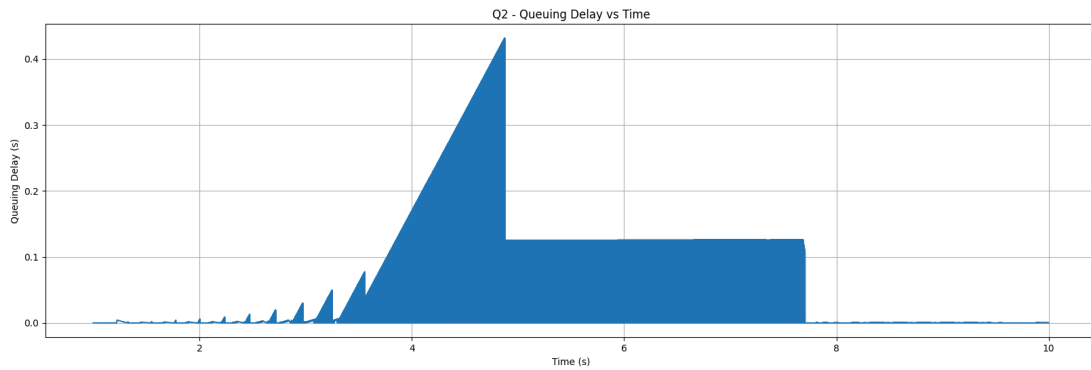
b) Plot CWND with time

- The plot will be saved as **Q2_cwnd.png** and will look like the following:

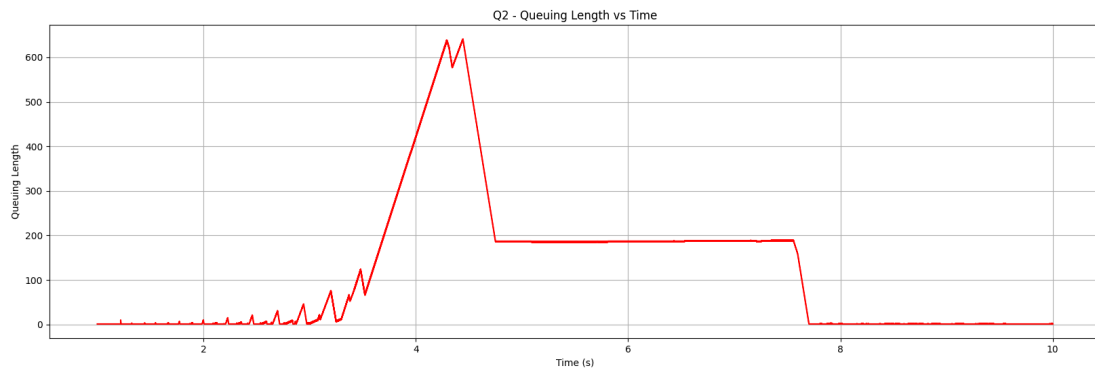


c) Plot queueing delay with time

- The plot will be saved as **Q2_QueueingDelay.png** and will look like the following:

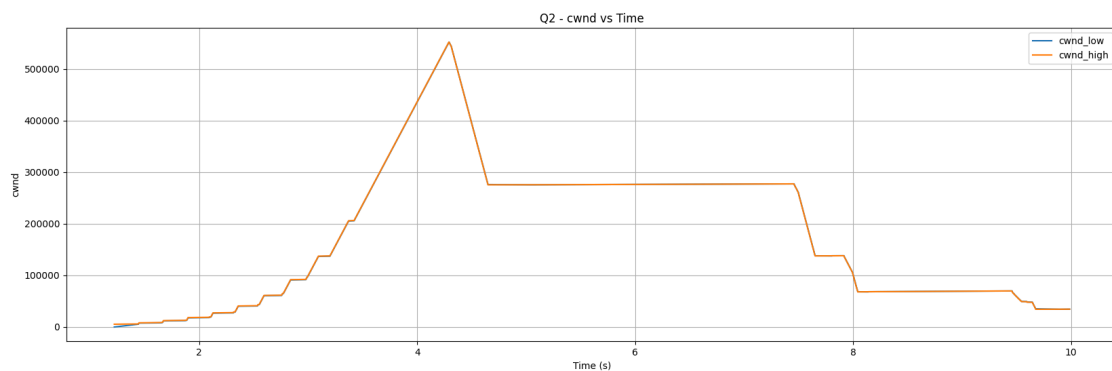



- The queueing length plot will be saved as **Q2_QueueingLength.png** and will look like the following:



d) Compare CWND plots of Q.1. and Q.2.; what insights did you gain?

- The CWND plot for Q1 and Q2 are as follows:



- Since the queue size is increased to **1000** packets in Q2, the congestion window size is larger in Q2 compared to Q1. The larger congestion window size allows more packets to be sent in the network, which increases the throughput of the network. Therefore,

the throughput of the network is higher in Q2 compared to Q1.

Q3 - Change N1-N2 bandwidth to 10 Mbps and N1-N2 delay to 100ms (rest of the parameter values are same as default values)

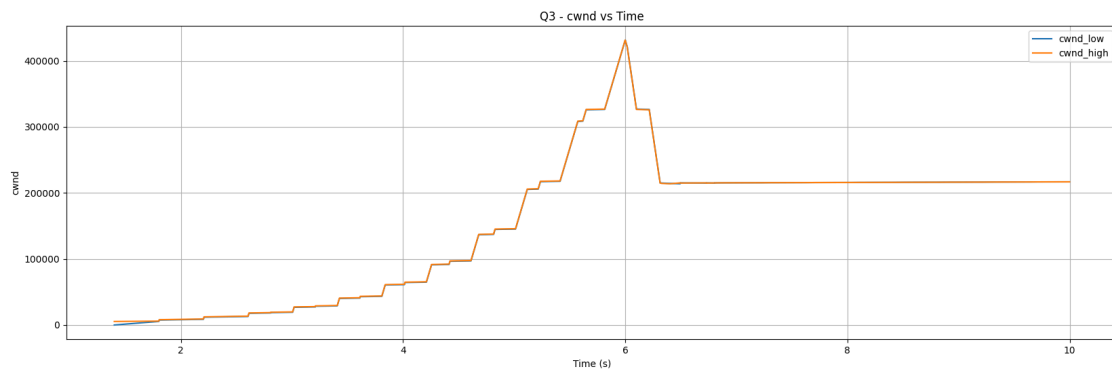
a) What is the average computed throughput of the TCP transfer?

- The average throughput values are:

```
Throughput for Q3_tcp-example-0-0.pcap is 3.419889777777778 Mbps  
Throughput for Q3_tcp-example-1-0.pcap is 3.378456006292842 Mbps  
Throughput for Q3_tcp-example-1-1.pcap is 3.378456006292842 Mbps  
Throughput for Q3_tcp-example-2-0.pcap is 3.4509745237548772 Mbps
```

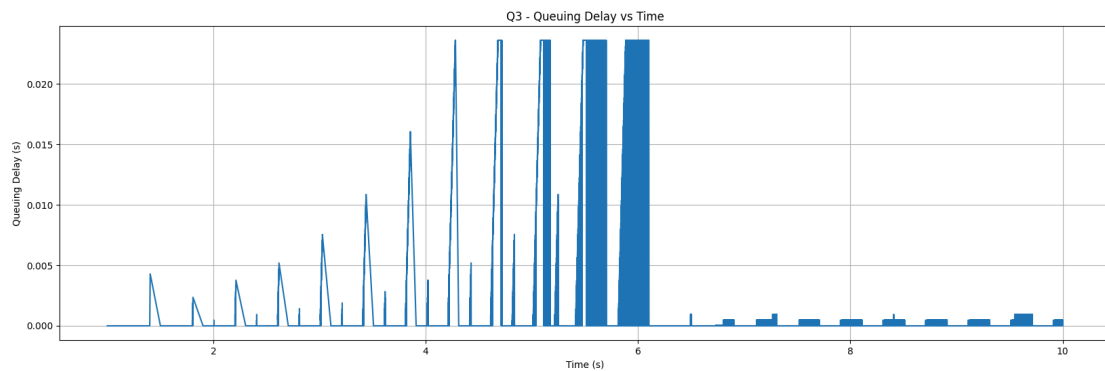
b) Plot CWND with time

- The plot will be saved as **Q3_cwnd.png** and will look like the following:

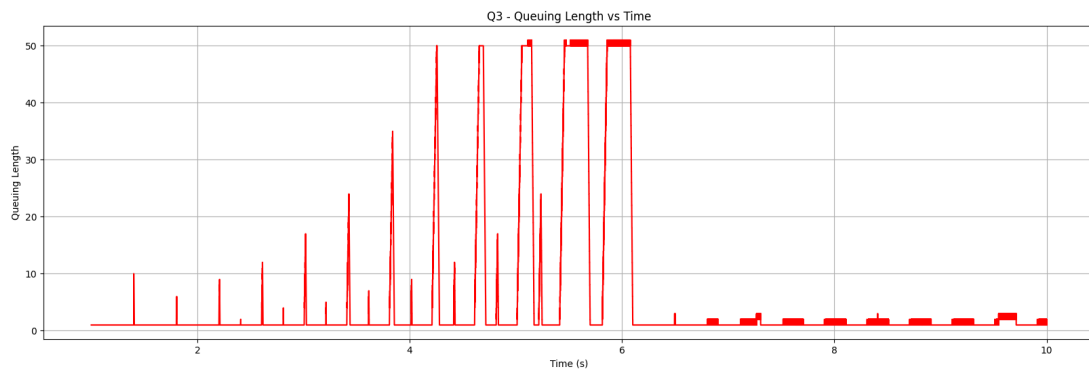


c) Plot queueing delay with time

- The plot will be saved as **Q3_QueueingDelay.png** and will look like the following:

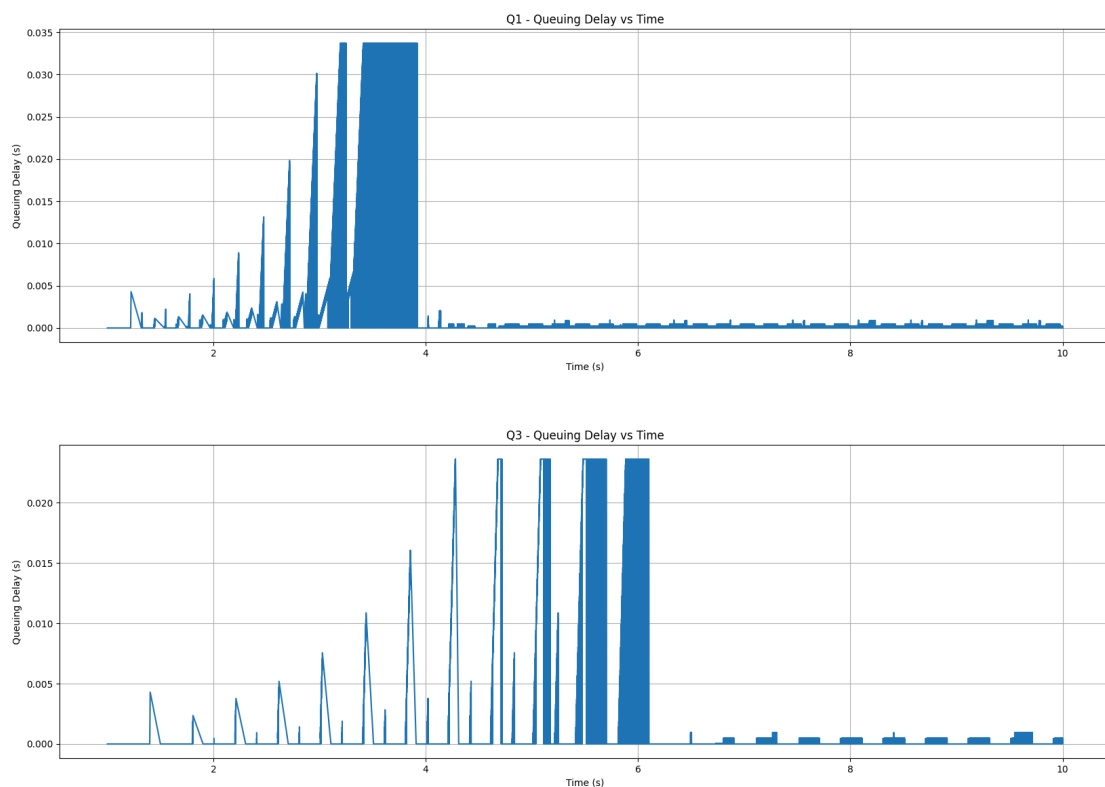


- The queueing length plot will be saved as **Q3_QueueingLength.png** and will look like the following:



d) Compare queueing delay plots of Q.1. and Q.3.; what insights did you gain?

- The queueing delay plot for Q1 and Q3 are as follows:



- The queueing delay is higher in Q1 compared to Q3 as the packets are sent more slowly in Q1 due to the lower bandwidth of the link between node **n1** and **n2**. The higher bandwidth of the second link in Q3 allows more packets to be sent in the network, which reduces the queueing delay in the network.
- We also see that the queueing delay takes longer to saturate in Q3 compared to Q1, which is due to the higher bandwidth of the link between node **n1** and **n2**.